

**Objective :** To implement a Multi-Layer Perceptron (MLP) using NumPy that can learn the XOR Boolean function.

**Description of the model :** The MLP consists of an input layer with 2 neurons, a hidden layer with 2 neurons, and an output layer with 1 neuron.

The activation function used is the **sigmoid function**.

The network is trained using backpropagation with **gradient descent**.

## Python Implementation :

```
import numpy as np

def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def sigmoid_derivative(x):
    return x * (1 - x)

X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
y = np.array([[0], [1], [1], [0]])

input_neurons = 2
hidden_neurons = 2
output_neurons = 1

np.random.seed(42)

weights_input_hidden = np.random.uniform(-1, 1, (input_neurons, hidden_neurons))
weights_hidden_output = np.random.uniform(-1, 1, (hidden_neurons, output_neurons))
bias_hidden = np.random.uniform(-1, 1, (1, hidden_neurons))
bias_output = np.random.uniform(-1, 1, (1, output_neurons))

learning_rate = 0.5
```

```
epochs = 10000
```

```
for epoch in range(epochs):
```

```
    hidden_layer_input = np.dot(X, weights_input_hidden) + bias_hidden
```

```
    hidden_layer_output = sigmoid(hidden_layer_input)
```

```
    output_layer_input = np.dot(hidden_layer_output, weights_hidden_output) + bias_output
```

```
    output = sigmoid(output_layer_input)
```

```
    error = y - output
```

```
    d_output = error * sigmoid_derivative(output)
```

```
    error_hidden_layer = d_output.dot(weights_hidden_output.T)
```

```
    d_hidden_layer = error_hidden_layer * sigmoid_derivative(hidden_layer_output)
```

```
    weights_hidden_output += hidden_layer_output.T.dot(d_output) * learning_rate
```

```
    bias_output += np.sum(d_output, axis=0, keepdims=True) * learning_rate
```

```
    weights_input_hidden += X.T.dot(d_hidden_layer) * learning_rate
```

```
    bias_hidden += np.sum(d_hidden_layer, axis=0, keepdims=True) * learning_rate
```

```
print("Final Output:")
```

```
print(output)
```

## Description of code :

- The model initializes random weights and biases.
- The forward pass computes activations for the hidden and output layers.
- The error is computed, and gradients are backpropagated to update weights.
- Training continues for 10,000 epochs with a learning rate of 0.5.

## Performance Evaluation :

The final output approximates the XOR truth table:

*Input: [0,0] -> Output  $\approx 0$*

*Input: [0,1] -> Output  $\approx 1$*

*Input: [1,0] -> Output  $\approx 1$*

*Input: [1,1] -> Output  $\approx 0$*

The network successfully learns the XOR function.